

This project demonstrates the application of Facebook AI's SAM2 model for a **zero-shot object detection task**. The primary objective was to use a single "setup" image and its corresponding mask for a given product from the CMU-10 3D dataset to automatically detect that product in subsequent images. The performance of these detections was then quantitatively measured for each product category using the Intersection over Union (IoU) metric

```
!pip install torch torchvision torchaudio  
!pip install Pillow numpy matplotlib  
!pip install pycocotools  
!pip install git+https://github.com/facebookresearch/segment-anything-2.git
```



```
# Download the CMU-10 3D Dataset
!wget http://www.cs.cmu.edu/~ehsiao/3drecognition/CMU10_3D.zip

# Unzip the dataset (-q makes the output quiet)
!unzip -q CMU10_3D.zip

→ --2025-08-07 04:13:01-- http://www.cs.cmu.edu/~ehsiao/3drecognition/CMU10\_3D.zip
Resolving www.cs.cmu.edu (www.cs.cmu.edu)... 128.2.42.95
Connecting to www.cs.cmu.edu (www.cs.cmu.edu).|128.2.42.95|:80... connected.
HTTP request sent, awaiting response... 200 OK
Length: 121503865 (116M) [application/zip]
Saving to: 'CMU10_3D.zip'

CMU10_3D.zip      100%[=====] 115.88M  1.98MB/s    in 72s

2025-08-07 04:14:14 (1.61 MB/s) - 'CMU10_3D.zip' saved [121503865/121503865]
```

Start coding or generate with AI.

```
import numpy as np
from PIL import Image
import matplotlib.pyplot as plt
import os

def process_img_png_mask(image_path, mask_path, visualize=False):
    """
    Processes a mask to extract the bounding box of the object.

    Args:
        image_path (str): Path to the full image.
        mask_path (str): Path to the mask image.
        visualize (bool): If True, displays the image, mask, and extracted object.

    Returns:
        list: A list containing [xmin, xmax, ymin, ymax].
    """
    mask_img = Image.open(mask_path).convert('L')
    mask = np.array(mask_img) > 0

    rows = np.any(mask, axis=1)
    cols = np.any(mask, axis=0)
    ymin, ymax = np.where(rows)[0][[0, -1]]
    xmin, xmax = np.where(cols)[0][[0, -1]]

    if visualize:
        full_img = Image.open(image_path)
        extracted_img = full_img.crop((xmin, ymin, xmax, ymax))

        fig, ax = plt.subplots(1, 3, figsize=(15, 5))
        ax[0].imshow(full_img)
        ax[0].set_title("Full Image")
        ax[1].imshow(mask_img, cmap='gray')
        ax[1].set_title("Mask")
        ax[2].imshow(extracted_img)
        ax[2].set_title("Extracted")
        plt.show()

    return [xmin, xmax, ymin, ymax]

!ls -l /content/
```

```
→ total 270928
drwxr-xr-x 5 root root      4096 May 20  2010 CMU10_3D
-rw-r--r-- 1 root root 121503865 May 20  2010 CMU10_3D.zip
-rw-r--r-- 1 root root 155906050 Aug  7 04:14 sam2_hiera_tiny.pt
drwxr-xr-x 1 root root      4096 Aug  5 17:22 sample_data

# Install the library to download from Hugging Face
!pip install huggingface-hub -q

from huggingface_hub import hf_hub_download
import shutil

print("Downloading model from Hugging Face...")
model_path = hf_hub_download(
    repo_id="facebook/sam2-hiera-tiny",
    filename="sam2_hiera_tiny.pt"
)
shutil.copy(model_path, "/content/sam2_hiera_tiny.pt")

print("Model successfully downloaded!")

→ Downloading model from Hugging Face...
/usr/local/lib/python3.11/dist-packages/huggingface_hub/utils/_auth.py:94: UserWarning:
The secret `HF_TOKEN` does not exist in your Colab secrets.
To authenticate with the Hugging Face Hub, create a token in your settings tab (https://huggingface.co/settings/token)
You will be able to reuse this secret in all of your notebooks.
Please note that authentication is recommended but still optional to access public models or datasets.
  warnings.warn(
sam2_hiera_tiny.pt: 100%                                         156M/156M [00:02<00:00, 82.9MB/s]
Model successfully downloaded!

!ls /content/

→ CMU10_3D      __pycache__      sam2_hiera_tiny.pt
CMU10_3D.zip  sam2_extract.py  sample_data

# This code block writes the necessary functions into a file named sam2_extract.py.
sam2_code = """
import torch
from sam2.build_sam import build_sam2_video_predictor
from PIL import Image
import numpy as np
import matplotlib.pyplot as plt
import os, glob, shutil
import matplotlib.patches as patches

# --- Model and Predictor Initialization ---
# Use the FULL, absolute path to the model and config
checkpoint = "/content/sam2_hiera_tiny.pt"
model_cfg = "sam2_hiera_t.yaml"

# Add a check to ensure the model file exists before loading
if not os.path.exists(checkpoint):
    raise FileNotFoundError(f"CRITICAL ERROR: Model checkpoint not found at: {checkpoint}")

predictor_vid = build_sam2_video_predictor(model_cfg, checkpoint, device='cuda')
tempfolder = "./tempdir"

# --- Helper Functions (as provided in the assignment) ---
def create_if_not_exists(dirname):
    if not os.path.exists(dirname):
        os.mkdir(dirname)
"""

# Write the code to a file
with open('sam2_extract.py', 'w') as f:
    f.write(sam2_code)
```

```

def cleardir(tempfolder):
    if os.path.exists(tempfolder):
        filepaths = glob.glob(tempfolder + "/*")
        for filepath in filepaths:
            os.unlink(filepath)

def show_mask(mask, ax, obj_id=None):
    cmap = plt.get_cmap("tab10")
    cmap_idx = 0 if obj_id is None else obj_id
    color = np.array([*cmap(cmap_idx)[3], 0.6])
    h, w = mask.shape[-2:]
    mask_image = mask.reshape(h, w, 1) * color.reshape(1, 1, -1)
    ax.imshow(mask_image)

# --- Main Tracking Function (as provided in the assignment) ---
def track_item_boxes(imgpath1, imgpath2, img1boxclasslist, visualize=True):
    create_if_not_exists(tempfolder)
    cleardir(tempfolder)

    shutil.copy(imgpath1, tempfolder + "/00000.jpg")
    shutil.copy(imgpath2, tempfolder + "/00001.jpg")

    video_path = "./tempdir"
    inference_state = predictor_vid.init_state(video_path=video_path)

    ann_frame_idx = 0
    for img1boxclass in img1boxclasslist:
        ([xmin, xmax, ymin, ymax], objectnumint) = img1boxclass
        box = np.array([xmin, ymin, xmax, ymax], dtype=np.float32)
        predictor_vid.add_new_points_or_box(
            inference_state=inference_state,
            frame_idx=ann_frame_idx,
            obj_id=objectnumint,
            box=box,
        )

    video_segments = {}
    for out_frame_idx, out_obj_ids, out_mask_logits in predictor_vid.propagate_in_video(inference_state):
        video_segments[out_frame_idx] = {
            out_obj_id: (out_mask_logits[i] > 0.0).cpu().numpy()
            for i, out_obj_id in enumerate(out_obj_ids)
        }

    if visualize:
        ([xmin, xmax, ymin, ymax], _) = img1boxclasslist[0]
        fig, ax = plt.subplots()
        plt.title(f"Original Image Object")
        ax.imshow(Image.open(os.path.join(video_path, "00000.jpg")))
        rect = patches.Rectangle((xmin, ymin), xmax - xmin, ymax - ymin, linewidth=1, edgecolor='r', facecolor='none')
        ax.add_patch(rect)
        plt.show()

        out_frame_idx = 1
        plt.figure(figsize=(6, 4))
        plt.title(f"Detected Object in Test Image")
        plt.imshow(Image.open(os.path.join(video_path, "00001.jpg")))
        if out_frame_idx in video_segments:
            for out_obj_id, out_mask in video_segments[out_frame_idx].items():
                show_mask(out_mask, plt.gca(), obj_id=out_obj_id)
        plt.show()

    return video_segments
"""

```

```

# Write the code to the file
with open("sam2_extract.py", "w") as f:
    f.write(sam2_code)

import torch
from sam2.build_sam import build_sam
from sam2.automatic_mask_generator import SAM2AutomaticMaskGenerator
from sam2.sam2_image_predictor import SAM2ImagePredictor
from sam2.build_sam import build_sam2_video_predictor
from PIL import Image, ImageOps
import numpy as np
import matplotlib.pyplot as plt
import os, glob, shutil
import matplotlib.patches as patches

# --- Model and Predictor Initialization ---
checkpoint = "/content/sam2_hiera_tiny.pt"
model_cfg = "sam2_hiera_t.yaml"

# NOTE: The PDF has some duplicate/unclear initializations.
# The video predictor is the one used by the track_item_boxes function.
predictor_vid = build_sam2_video_predictor(model_cfg, checkpoint, device='cuda')
tempfolder = "./tempdir"

# --- Helper Functions ---
def create_if_not_exists(dirname):
    if not os.path.exists(dirname):
        os.mkdir(dirname)

def cleardir(tempfolder):
    filepaths = glob.glob(tempfolder + "/*")
    for filepath in filepaths:
        os.unlink(filepath)

def show_mask(mask, ax, obj_id=None, random_color=False):
    if random_color:
        color = np.concatenate([np.random.random(3), np.array([0.6])], axis=0)
    else:
        cmap = plt.get_cmap("tab10")
        cmap_idx = 0 if obj_id is None else obj_id
        color = np.array([*cmap(cmap_idx)[:3], 0.6])
    h, w = mask.shape[-2:]
    mask_image = mask.reshape(h, w, 1) * color.reshape(1, 1, -1)
    ax.imshow(mask_image)

def show_box(box, ax):
    x0, y0 = box[0], box[1]
    w, h = box[2] - box[0], box[3] - box[1]
    ax.add_patch(patches.Rectangle((x0, y0), w, h, edgecolor='green', facecolor=(0, 0, 0, 0), lw=2))

# --- Main Tracking Function ---
def track_item_boxes(imgpath1, imgpath2, img1boxclasslist, visualize=True):
    # imgpath1: Image where object is known
    # imgpath2: Image where object is to be tracked
    # img1boxclasslist: [[xmin, xmax, ymin, ymax], objectnumint,...] for all objects

    create_if_not_exists(tempfolder)
    cleardir(tempfolder)

    shutil.copy(imgpath1, tempfolder + "/00000.jpg")
    shutil.copy(imgpath2, tempfolder + "/00001.jpg")

    inference_state = predictor_vid.init_state(video_path="./tempdir")
    predictor_vid.reset_state(inference_state)

```

```

ann_frame_idx = 0
for img1boxclass in img1boxclasslist:
    ([xmin, xmax, ymin, ymax], objectnumint) = img1boxclass
    box = np.array([xmin, ymin, xmax, ymax], dtype=np.float32)
    out_obj_ids, out_mask_logits = predictor_vid.add_new_points_or_box(
        inference_state=inference_state,
        frame_idx=ann_frame_idx,
        obj_id=objectnumint,
        box=box,
    )

video_segments = {}
for out_frame_idx, out_obj_ids, out_mask_logits in predictor_vid.propagate_in_video(inference_state):
    video_segments[out_frame_idx] = {
        out_obj_id: (out_mask_logits[i] > 0.0).cpu().numpy()
        for i, out_obj_id in enumerate(out_obj_ids)
    }

if visualize:
    # Visualize the source image with its bounding box
    fig, ax = plt.subplots()
    plt.title(f"original image object ::")
    ax.imshow(Image.open(tempfolder + "/00000.jpg"))
    rect = patches.Rectangle((xmin, ymin), xmax - xmin, ymax - ymin, linewidth=1, edgecolor='r', facecolor='none')
    ax.add_patch(rect)
    plt.show()

    # Visualize the target image with the predicted mask
    out_frame_idx = 1
    plt.figure(figsize=(6, 4))
    plt.title(f"detected object in test image ::")
    plt.imshow(Image.open(tempfolder + "/00001.jpg"))
    if out_frame_idx in video_segments:
        for out_obj_id, out_mask in video_segments[out_frame_idx].items():
            show_mask(out_mask, plt.gca(), obj_id=out_obj_id)
    plt.show()

return video_segments

```

!ls /content/CMU10_3D/

→ data_2D data_3D modelimages README.txt

```

from sam2_extract import track_item_boxes

# Path to the dataset
dataset_path = '/content/CMU10_3D/data_2D/'
product_folders = [f for f in os.listdir(dataset_path) if os.path.isdir(os.path.join(dataset_path, f))]

for product in product_folders:
    product_path = os.path.join(dataset_path, product)
    images = sorted([os.path.join(product_path, img) for img in os.listdir(product_path) if img.endswith('.jpg')])
    masks = sorted([os.path.join(product_path, msk) for msk in os.listdir(product_path) if msk.endswith('.png')])

    first_img_path = images[0]
    first_img_mask_path = masks[0]

    # Get the bounding box from the first image's mask
    [xmin, xmax, ymin, ymax] = process_img_png_mask(first_img_path, first_img_mask_path)

    # Detect the object in the rest of the images
    for i in range(1, len(images)):
        test_img_path = images[i]

```

```

# The '1' is a dummy object ID as mentioned in the assignment
output = track_item_boxes(first_img_path, test_img_path, [[xmin, xmax, ymin, ymax], 1]), visualize=True)

# The output mask for the test image
output_masks = output[1]
relevant_mask = output_masks[1]

# Further steps will involve converting this mask to a bounding box and calculating accuracy

def get_bbox_from_mask(mask):
    """Converts a boolean mask to a bounding box."""
    rows = np.any(mask, axis=1)
    cols = np.any(mask, axis=0)
    if not np.any(rows) or not np.any(cols):
        return [0, 0, 0, 0]
    ymin, ymax = np.where(rows)[0][[0, -1]]
    xmin, xmax = np.where(cols)[0][[0, -1]]
    return [xmin, xmax, ymin, ymax]

def calculate_iou(boxA, boxB):
    """Calculates Intersection over Union for two bounding boxes."""
    xA = max(boxA[0], boxB[0])
    yA = max(boxA[2], boxB[2])
    xB = min(boxA[1], boxB[1])
    yB = min(boxA[3], boxB[3])

    interArea = max(0, xB - xA + 1) * max(0, yB - yA + 1)

    boxAArea = (boxA[1] - boxA[0] + 1) * (boxA[3] - boxA[2] + 1)
    boxBArea = (boxB[1] - boxB[0] + 1) * (boxB[3] - boxB[2] + 1)

    iou = interArea / float(boxAArea + boxBArea - interArea)
    return iou

# --- Inside the loop for each product ---
# (Code from step 3 continues here)

total_iou = 0
num_images = 0

for i in range(1, len(images)):
    test_img_path = images[i]
    ground_truth_mask_path = masks[i]

    output = track_item_boxes(first_img_path, test_img_path, [[xmin, xmax, ymin, ymax], 1]), visualize=False)

    if 1 in output and 1 in output[1]:
        predicted_mask = output[1][1]
        predicted_box = get_bbox_from_mask(predicted_mask)

        ground_truth_box = process_img_png_mask(test_img_path, ground_truth_mask_path)

        iou = calculate_iou(predicted_box, ground_truth_box)
        total_iou += iou
        num_images += 1
        print(f"  IoU for {os.path.basename(test_img_path)}: {iou:.4f}")

if num_images > 0:
    average_iou = total_iou / num_images
    print(f"\nAverage IoU for product '{product}': {average_iou:.4f}\n")

import os
import numpy as np

```

```

from PIL import Image
import matplotlib.pyplot as plt
import time # Import time to see how long it takes

# --- This should be the only import you need now ---
from sam2_extract import track_item_boxes

print("All libraries imported successfully.")
print("-" * 50)

# --- Helper Function 1: Get Bounding Box from a Ground Truth Mask ---
def process_img_png_mask(image_path, mask_path):
    try:
        with Image.open(mask_path) as mask_img:
            mask = np.array(mask_img.convert('L')) > 0
            rows, cols = np.any(mask, axis=1), np.any(mask, axis=0)
            if not np.any(rows) or not np.any(cols): return None
            ymin, ymax = np.where(rows)[0][[0, -1]]
            xmin, xmax = np.where(cols)[0][[0, -1]]
            return [xmin, xmax, ymin, ymax]
    except Exception as e:
        print(f"  ERROR reading mask {mask_path}: {e}")
        return None

# --- Helper Function 2: Convert a Predicted Boolean Mask to a Bounding Box ---
def get_bbox_from_mask(mask):
    rows, cols = np.any(mask, axis=1), np.any(mask, axis=0)
    if not np.any(rows) or not np.any(cols): return [0, 0, 0, 0]
    ymin, ymax = np.where(rows)[0][[0, -1]]
    xmin, xmax = np.where(cols)[0][[0, -1]]
    return [xmin, xmax, ymin, ymax]

# --- Helper Function 3: Calculate Intersection over Union (IoU) ---
def calculate_iou(boxA, boxB):
    xA = max(boxA[0], boxB[0]); yA = max(boxA[2], boxB[2])
    xB = min(boxA[1], boxB[1]); yB = min(boxA[3], boxB[3])
    interArea = max(0, xB - xA + 1) * max(0, yB - yA + 1)
    boxAArea = (boxA[1] - boxA[0] + 1) * (boxA[3] - boxA[2] + 1)
    boxBArea = (boxB[1] - boxB[0] + 1) * (boxB[3] - boxB[2] + 1)
    iou = interArea / float(boxAArea + boxBArea - interArea)
    return iou

# --- Main Execution ---
print("Starting the main object tracking script...")

dataset_path = '/content/CMU10_3D/data_2D/'
if not os.path.exists(dataset_path):
    print(f"FATAL ERROR: The dataset path does not exist: {dataset_path}")
else:
    product_folders = sorted([f for f in os.listdir(dataset_path) if os.path.isdir(os.path.join(dataset_path, f))])
    print(f"Found {len(product_folders)} product folders to process.")
    print("-" * 50)

    # Loop through each product folder
    for product in product_folders:
        start_time = time.time()
        print(f" Processing Product: {product}")
        product_path = os.path.join(dataset_path, product)

        images = sorted([os.path.join(product_path, img) for img in os.listdir(product_path) if img.endswith('.jpg')])
        masks = sorted([os.path.join(product_path, msk) for msk in os.listdir(product_path) if msk.endswith('.png')])

        if not images or not masks:
            print("  Skipping: No images or masks found.")
            continue

```

```

first_img_path = images[0]
first_img_mask_path = masks[0]
initial_bbox = process_img_png_mask(first_img_path, first_img_mask_path)

if initial_bbox is None:
    print(f" Skipping product {product}: No valid object found in the first mask.")
    continue

iou_scores = []
# Loop through the rest of the images
for i in range(1, len(images)):
    test_img_path = images[i]
    ground_truth_mask_path = masks[i]

    # Set visualize=False to run faster
    video_segments = track_item_boxes(first_img_path, test_img_path, [(initial_bbox, 1)], visualize=False)

    predicted_box = [0, 0, 0, 0]
    if 1 in video_segments and 1 in video_segments[1]:
        predicted_mask = video_segments[1][1]
        predicted_box = get_bbox_from_mask(predicted_mask)

    ground_truth_box = process_img_png_mask(test_img_path, ground_truth_mask_path)

    if ground_truth_box:
        iou = calculate_iou(predicted_box, ground_truth_box)
        iou_scores.append(iou)
        # This is the line that prints the result for each image
        print(f"      Image: {os.path.basename(test_img_path)} | IoU Score: {iou:.4f}")

# Calculate and print the average for the product
if iou_scores:
    average_iou = sum(iou_scores) / len(iou_scores)
    end_time = time.time()
    print("\n RESULTS for '{product}':")
    print(f"      Average IoU: {average_iou:.4f}")
    print(f"      Time Taken: {end_time - start_time:.2f} seconds")
else:
    print("      No test images were processed for this product.")
print("-" * 50)

```

→ All libraries imported successfully.

Starting the main object tracking script...
Found 0 product folders to process.

```
# List the contents of the unzipped directory
!ls /content/CMU10_3D/
```

→ data_2D data_3D modelimages README.txt

```

import os
import numpy as np
from PIL import Image
import matplotlib.pyplot as plt
import time

# Import the main function from the file we created
from sam2_extract import track_item_boxes

print(" All libraries imported successfully.")
print("-" * 50)

```

```

# --- Helper functions (no changes here) ---
def process_img_png_mask(image_path, mask_path):
    try:
        with Image.open(mask_path) as mask_img:
            mask = np.array(mask_img.convert('L')) > 0
            rows, cols = np.any(mask, axis=1), np.any(mask, axis=0)
            if not np.any(rows) or not np.any(cols): return None
            ymin, ymax = np.where(rows)[0][[0, -1]]
            xmin, xmax = np.where(cols)[0][[0, -1]]
            return [xmin, xmax, ymin, ymax]
    except Exception as e:
        print(f"  ERROR reading mask {mask_path}: {e}")
        return None

def get_bbox_from_mask(mask):
    rows, cols = np.any(mask, axis=1), np.any(mask, axis=0)
    if not np.any(rows) or not np.any(cols): return [0, 0, 0, 0]
    ymin, ymax = np.where(rows)[0][[0, -1]]
    xmin, xmax = np.where(cols)[0][[0, -1]]
    return [xmin, xmax, ymin, ymax]

def calculate_iou(boxA, boxB):
    xA = max(boxA[0], boxB[0]); yA = max(boxA[2], boxB[2])
    xB = min(boxA[1], boxB[1]); yB = min(boxA[3], boxB[3])
    interArea = max(0, xB - xA + 1) * max(0, yB - yA + 1)
    boxAArea = (boxA[1] - boxA[0] + 1) * (boxA[3] - boxA[2] + 1)
    boxBArea = (boxB[1] - boxB[0] + 1) * (boxB[3] - boxB[2] + 1)
    iou = interArea / float(boxAArea + boxBArea - interArea)
    return iou

# --- Main Execution ---
print(" Starting the main object tracking script...")

# THIS IS THE CONFIRMED CORRECT PATH
dataset_path = '/content/CMU10_3D/data_2D/'

if not os.path.exists(dataset_path):
    print(f" FATAL ERROR: The dataset path still does not exist: {dataset_path}")
else:
    product_folders = sorted([f for f in os.listdir(dataset_path) if os.path.isdir(os.path.join(dataset_path, f))])
    print(f"Found {len(product_folders)} product folders to process.")
    print("-" * 50)

    for product in product_folders:
        start_time = time.time()
        print(f" Processing Product: {product}")
        product_path = os.path.join(dataset_path, product)

        images = sorted([os.path.join(product_path, img) for img in os.listdir(product_path) if img.endswith('.jpg')])
        masks = sorted([os.path.join(product_path, msk) for msk in os.listdir(product_path) if msk.endswith('.png')])

        if not images or not masks:
            print("  Skipping: No images or masks found.")
            continue

        first_img_path = images[0]
        first_img_mask_path = masks[0]
        initial_bbox = process_img_png_mask(first_img_path, first_img_mask_path)

        if initial_bbox is None:
            print(f"  Skipping product {product}: No valid object found in the first mask.")
            continue

        iou_scores = []

```

```

for i in range(1, len(images)):
    test_img_path = images[i]
    ground_truth_mask_path = masks[i]

    video_segments = track_item_boxes(first_img_path, test_img_path, [(initial_bbox, 1)], visualize=False)

    predicted_box = [0, 0, 0, 0]
    if 1 in video_segments and 1 in video_segments[1]:
        predicted_mask = video_segments[1][1]
        predicted_box = get_bbox_from_mask(predicted_mask)

    ground_truth_box = process_img_png_mask(test_img_path, ground_truth_mask_path)

    if ground_truth_box:
        iou = calculate_iou(predicted_box, ground_truth_box)
        iou_scores.append(iou)
        print(f"    Image: {os.path.basename(test_img_path)} | IoU Score: {iou:.4f}")

if iou_scores:
    average_iou = sum(iou_scores) / len(iou_scores)
    end_time = time.time()
    print(f"\nRESULTS for '{product}'")
    print(f"    Average IoU: {average_iou:.4f}")
    print(f"    Time Taken: {end_time - start_time:.2f} seconds")
else:
    print("    No test images were processed for this product.")
print("-" * 50)

```

→ All libraries imported successfully.

Starting the main object tracking script...
Found 0 product folders to process.

```

import os
import numpy as np
from PIL import Image
import matplotlib.pyplot as plt
import time
from collections import defaultdict

# --- Step 1: Ensure sam2_extract.py is created correctly ---
# This code block writes the necessary functions into a file named sam2_extract.py.
# This ensures it's in the correct directory for your notebook to import.
sam2_code = """
import torch
from sam2.build_sam import build_sam2_video_predictor
from PIL import Image
import numpy as np
import matplotlib.pyplot as plt
import os, glob, shutil
import matplotlib.patches as patches

# --- Model and Predictor Initialization ---
checkpoint = "/content/sam2_hiera_tiny.pt"
model_cfg = "sam2_hiera_t.yaml"
predictor_vid = build_sam2_video_predictor(model_cfg, checkpoint, device='cuda')
tempfolder = "./tempdir"

def create_if_not_exists(dirname):
    if not os.path.exists(dirname):
        os.mkdir(dirname)

def cleardir(tempfolder):
    if os.path.exists(tempfolder):

```

```

filepaths = glob.glob(tempfolder + "/*")
for filepath in filepaths:
    os.unlink(filepath)

def show_mask(mask, ax, obj_id=None):
    cmap = plt.get_cmap("tab10")
    cmap_idx = 0 if obj_id is None else obj_id
    color = np.array([*cmap(cmap_idx)[:3], 0.6])
    h, w = mask.shape[-2:]
    mask_image = mask.reshape(h, w, 1) * color.reshape(1, 1, -1)
    ax.imshow(mask_image)

def track_item_boxes(imgpath1, imgpath2, img1boxclasslist, visualize=True):
    create_if_not_exists(tempfolder)
    cleardir(tempfolder)
    shutil.copy(imgpath1, os.path.join(tempfolder, "00000.jpg"))
    shutil.copy(imgpath2, os.path.join(tempfolder, "00001.jpg"))
    inference_state = predictor_vid.init_state(video_path="./tempdir")
    ann_frame_idx = 0
    for img1boxclass in img1boxclasslist:
        ([xmin, xmax, ymin, ymax], objectnumint) = img1boxclass
        box = np.array([xmin, ymin, xmax, ymax], dtype=np.float32)
        predictor_vid.add_new_points_or_box(
            inference_state=inference_state,
            frame_idx=ann_frame_idx,
            obj_id=objectnumint,
            box=box,
        )
    video_segments = {}
    for out_frame_idx, out_obj_ids, out_mask_logits in predictor_vid.propagate_in_video(inference_state):
        video_segments[out_frame_idx] = {
            out_obj_id: (out_mask_logits[i] > 0.0).cpu().numpy()
            for i, out_obj_id in enumerate(out_obj_ids)
        }
    if visualize:
        ([xmin, xmax, ymin, ymax], _) = img1boxclasslist[0]
        fig, ax = plt.subplots()
        plt.title(f"Original Image Object")
        ax.imshow(Image.open(os.path.join(tempfolder, "00000.jpg")))
        rect = patches.Rectangle((xmin, ymin), xmax - xmin, ymax - ymin, linewidth=1, edgecolor='r', facecolor='none')
        ax.add_patch(rect)
        plt.show()
        out_frame_idx = 1
        plt.figure(figsize=(6, 4))
        plt.title(f"Detected Object in Test Image")
        plt.imshow(Image.open(os.path.join(tempfolder, "00001.jpg")))
        if out_frame_idx in video_segments:
            for out_obj_id, out_mask in video_segments[out_frame_idx].items():
                show_mask(out_mask, plt.gca(), obj_id=out_obj_id)
        plt.show()
    return video_segments
"""

with open("sam2_extract.py", "w") as f:
    f.write(sam2_code)
print(`sam2_extract.py` created successfully.)

# --- Step 2: Main processing script ---
from sam2_extract import track_item_boxes

# Helper functions for bounding box processing and IoU calculation
def process_img_png_mask(image_path, mask_path):
    try:
        with Image.open(mask_path) as mask_img:
            mask = np.array(mask_img.convert('L')) > 0
            rows, cols = np.any(mask, axis=1), np.any(mask, axis=0)

```

```

if not np.any(rows) or not np.any(cols): return None
ymin, ymax = np.where(rows)[0][[0, -1]]
xmin, xmax = np.where(cols)[0][[0, -1]]
return [xmin, xmax, ymin, ymax]
except (FileNotFoundException, Exception) as e:
    return None

def get_bbox_from_mask(mask):
    rows, cols = np.any(mask, axis=1), np.any(mask, axis=0)
    if not np.any(rows) or not np.any(cols): return [0, 0, 0, 0]
    ymin, ymax = np.where(rows)[0][[0, -1]]
    xmin, xmax = np.where(cols)[0][[0, -1]]
    return [xmin, xmax, ymin, ymax]

def calculate_iou(boxA, boxB):
    xA = max(boxA[0], boxB[0]); yA = max(boxA[2], boxB[2])
    xB = min(boxA[1], boxB[1]); yB = min(boxA[3], boxB[3])
    interArea = max(0, xB - xA + 1) * max(0, yB - yA + 1)
    boxAArea = (boxA[1] - boxA[0] + 1) * (boxA[3] - boxA[2] + 1)
    boxBArea = (boxB[1] - boxB[0] + 1) * (boxB[3] - boxB[2] + 1)
    iou = interArea / float(boxAArea + boxBArea - interArea)
    return iou

# Main Execution with corrected file discovery logic
print("Starting the main object tracking script...")
dataset_path = '/content/CMU10_3D/data_2D/'

if os.path.exists(dataset_path):
    all_files = os.listdir(dataset_path)
    product_files = defaultdict(lambda: {'images': [], 'masks': []})

    for f in all_files:
        if f.endswith('.jpg'):
            product_name = '_'.join(f.split('_')[:-1])
            product_files[product_name]['images'].append(f)
        elif f.endswith('.png'):
            product_name = '_'.join(f.split('_')[:-3])
            product_files[product_name]['masks'].append(f)

    print(f"Found {len(product_files)} product types to process.")
    print("-" * 50)

    for product, files in product_files.items():
        start_time = time.time()
        print(f" Processing Product: {product}")

        images = sorted([os.path.join(dataset_path, img) for img in files['images']])
        masks = sorted([os.path.join(dataset_path, msk) for msk in files['masks']])

        if not images or not masks or len(images) != len(masks) or len(images) < 2:
            print(f" Skipping {product}: Not enough image/mask pairs for tracking.")
            continue

        first_img_path = images[0]
        first_img_mask_path = masks[0]
        initial_bbox = process_img_png_mask(first_img_path, first_img_mask_path)

        if initial_bbox is None:
            print(f" Skipping {product}: Could not find a valid bounding box in the first mask.")
            continue

        iou_scores = []
        for i in range(1, len(images)):
            test_img_path = images[i]
            ground_truth_mask_path = masks[i]

```

```

video_segments = track_item_boxes(first_img_path, test_img_path, [(initial_bbox, 1)], visualize=False)

predicted_box = [0, 0, 0, 0]
if 1 in video_segments and 1 in video_segments[1]:
    predicted_mask = video_segments[1][1]
    predicted_box = get_bbox_from_mask(predicted_mask)

ground_truth_box = process_img_png_mask(test_img_path, ground_truth_mask_path)

if ground_truth_box:
    iou = calculate_iou(predicted_box, ground_truth_box)
    iou_scores.append(iou)
    print(f"    Image: {os.path.basename(test_img_path)} | IoU Score: {iou:.4f}")

if iou_scores:
    average_iou = sum(iou_scores) / len(iou_scores)
    end_time = time.time()
    print(f"\n RESULTS for '{product}':")
    print(f"    Average IoU: {average_iou:.4f}")
    print(f"    Time Taken: {end_time - start_time:.2f} seconds")
else:
    print(f"    No valid comparisons were made for product '{product}'.")
print("-" * 50)
else:
    print(f"FATAL ERROR: The dataset path does not exist: {dataset_path}")

```

→ `sam2_extract.py` created successfully.
Starting the main object tracking script...
Found 10 product types to process.

Processing Product: can_chowder
Skipping can_chowder: Not enough image/mask pairs for tracking.

Processing Product: juicebox
Skipping juicebox: Not enough image/mask pairs for tracking.

Processing Product: ricepilaf
Skipping ricepilaf: Not enough image/mask pairs for tracking.

Processing Product: hc_potroastsoup
Skipping hc_potroastsoup: Not enough image/mask pairs for tracking.

Processing Product: rice_tuscan
Skipping rice_tuscan: Not enough image/mask pairs for tracking.

Processing Product: carton_oj
Skipping carton_oj: Not enough image/mask pairs for tracking.

Processing Product: can_soymilk
Skipping can_soymilk: Not enough image/mask pairs for tracking.

Processing Product: can_tomatosoup
Skipping can_tomatosoup: Not enough image/mask pairs for tracking.

Processing Product: diet_coke
Skipping diet_coke: Not enough image/mask pairs for tracking.

Processing Product: carton_soymilk
Skipping carton_soymilk: Not enough image/mask pairs for tracking.

```
!pip install git+https://github.com/facebookresearch/segment-anything-2.git -q
!pip install pycocotools -q
```

```

import os
import numpy as np
from PIL import Image
import matplotlib.pyplot as plt
import time
from collections import defaultdict

# --- Step 2: Ensure sam2_extract.py is created correctly ---
sam2_code = """
import torch
from sam2.build_sam import build_sam2_video_predictor
"""

```

```
from PIL import Image
import numpy as np
import matplotlib.pyplot as plt
import os, glob, shutil
import matplotlib.patches as patches
# --- Model and Predictor Initialization ---
checkpoint = "/content/sam2_hiera_tiny.pt"
model_cfg = "sam2_hiera_t.yaml"
# Add a check to ensure the model file exists before loading
if not os.path.exists(checkpoint):
    raise FileNotFoundError(f"CRITICAL ERROR: Model checkpoint not found at: {checkpoint}")
predictor_vid = build_sam2_video_predictor(model_cfg, checkpoint, device='cuda')
tempfolder = "./tempdir"
# --- Helper Functions (as provided in the assignment) ---
def create_if_not_exists(dirname):
    if not os.path.exists(dirname):
        os.mkdir(dirname)
def cleardir(tempfolder):
    if os.path.exists(tempfolder):
        filepaths = glob.glob(tempfolder + "/*")
        for filepath in filepaths:
            os.unlink(filepath)
def show_mask(mask, ax, obj_id=None):
    cmap = plt.get_cmap("tab10")
    cmap_idx = 0 if obj_id is None else obj_id
    color = np.array([*cmap(cmap_idx)[:3], 0.6])
    h, w = mask.shape[-2:]
    mask_image = mask.reshape(h, w, 1) * color.reshape(1, 1, -1)
    ax.imshow(mask_image)
# --- Main Tracking Function (as provided in the assignment) ---
def track_item_boxes(imgpath1, imgpath2, img1boxclasslist, visualize=True):
    create_if_not_exists(tempfolder)
    cleardir(tempfolder)
    shutil.copy(imgpath1, os.path.join(tempfolder, "00000.jpg"))
    shutil.copy(imgpath2, os.path.join(tempfolder, "00001.jpg"))
    inference_state = predictor_vid.init_state(video_path="./tempdir")
    ann_frame_idx = 0
    for img1boxclass in img1boxclasslist:
        ([xmin, xmax, ymin, ymax], objectnumint) = img1boxclass
        box = np.array([xmin, ymin, xmax, ymax], dtype=np.float32)
        predictor_vid.add_new_points_or_box(
            inference_state=inference_state,
            frame_idx=ann_frame_idx,
            obj_id=objectnumint,
            box=box,
        )
    video_segments = {}
    for out_frame_idx, out_obj_ids, out_mask_logits in predictor_vid.propagate_in_video(inference_state):
        video_segments[out_frame_idx] = {
            out_obj_id: (out_mask_logits[i] > 0.0).cpu().numpy()
            for i, out_obj_id in enumerate(out_obj_ids)
        }
    if visualize:
        ([xmin, xmax, ymin, ymax], _) = img1boxclasslist[0]
        fig, ax = plt.subplots()
        plt.title(f"Original Image Object")
        ax.imshow(Image.open(os.path.join(tempfolder, "00000.jpg")))
        rect = patches.Rectangle((xmin, ymin), xmax - xmin, ymax - ymin, linewidth=1, edgecolor='r', facecolor='none')
        ax.add_patch(rect)
        plt.show()
        out_frame_idx = 1
        plt.figure(figsize=(6, 4))
        plt.title(f"Detected Object in Test Image")
        plt.imshow(Image.open(os.path.join(tempfolder, "00001.jpg")))
        if out_frame_idx in video_segments:
            for out_obj_id, out_mask in video_segments[out_frame_idx].items():
                if out_obj_id in out_obj_ids:
```

```

        show_mask(out_mask, plt.gca(), obj_id=out_obj_id)
    plt.show()
    return video_segments
"""

with open("sam2_extract.py", "w") as f:
    f.write(sam2_code)
print(`sam2_extract.py` created successfully.)
from sam2_extract import track_item_boxes

# --- Step 3: Main Processing script with Debugging ---

# Helper functions for bounding box processing and IoU calculation
def process_img_png_mask(image_path, mask_path):
    try:
        with Image.open(mask_path) as mask_img:
            mask = np.array(mask_img.convert('L')) > 0
            rows, cols = np.any(mask, axis=1), np.any(mask, axis=0)
            if not np.any(rows) or not np.any(cols): return None
            ymin, ymax = np.where(rows)[0][[0, -1]]
            xmin, xmax = np.where(cols)[0][[0, -1]]
            return [xmin, xmax, ymin, ymax]
    except (FileNotFoundException, Exception):
        return None

def get_bbox_from_mask(mask):
    rows, cols = np.any(mask, axis=1), np.any(mask, axis=0)
    if not np.any(rows) or not np.any(cols): return [0, 0, 0, 0]
    ymin, ymax = np.where(rows)[0][[0, -1]]
    xmin, xmax = np.where(cols)[0][[0, -1]]
    return [xmin, xmax, ymin, ymax]

def calculate_iou(boxA, boxB):
    xA = max(boxA[0], boxB[0]); yA = max(boxA[2], boxB[2])
    xB = min(boxA[1], boxB[1]); yB = min(boxA[3], boxB[3])
    interArea = max(0, xB - xA + 1) * max(0, yB - yA + 1)
    boxAArea = (boxA[1] - boxA[0] + 1) * (boxA[3] - boxA[2] + 1)
    boxBArea = (boxB[1] - boxB[0] + 1) * (boxB[3] - boxB[2] + 1)
    iou = interArea / float(boxAArea + boxBArea - interArea)
    return iou

print("Starting the main object tracking script...")
dataset_path = '/content/CMU10_3D/data_2D/'

if os.path.exists(dataset_path):
    all_files = os.listdir(dataset_path)
    product_files = defaultdict(lambda: {'images': [], 'masks': []})

    for f in all_files:
        if f.endswith('.jpg'):
            product_name = '_'.join(f.split('_')[:-1])
            product_files[product_name]['images'].append(f)
        elif f.endswith('.png'):
            product_name = '_'.join(f.split('_')[:-3])
            product_files[product_name]['masks'].append(f)

    # --- DEBUGGING TABLE ---
    print("--- File Discovery Debug Table ---")
    print(f'{len(product_files)} products found')
    print(f'{sum(len(files["images"]) for _, files in product_files.items())} images found')
    print(f'{sum(len(files["masks"]) for _, files in product_files.items())} masks found')

    # --- Main Loop ---
    for product, files in product_files.items():
        start_time = time.time()

```

```

print(f"Processing Product: {product}")

images = sorted([os.path.join(dataset_path, img) for img in files['images']])
masks = sorted([os.path.join(dataset_path, msk) for msk in files['masks']])

if len(images) != len(masks) or len(images) < 2:
    print(f"  Skipping '{product}': Mismatch or not enough files (found {len(images)} images, {len(masks)} masks")
    print("-" * 50)
    continue

first_img_path, first_img_mask_path = images[0], masks[0]
initial_bbox = process_img_png_mask(first_img_path, first_img_mask_path)
if initial_bbox is None:
    print(f"  Skipping '{product}': Could not find a valid bounding box in the first mask.")
    print("-" * 50)
    continue

iou_scores = []
for i in range(1, len(images)):
    test_img_path, ground_truth_mask_path = images[i], masks[i]
    video_segments = track_item_boxes(first_img_path, test_img_path, [(initial_bbox, 1)], visualize=False)
    predicted_box = [0, 0, 0, 0]
    if 1 in video_segments and 1 in video_segments[1]:
        predicted_mask = video_segments[1][1]
        predicted_box = get_bbox_from_mask(predicted_mask)

    ground_truth_box = process_img_png_mask(test_img_path, ground_truth_mask_path)
    if ground_truth_box:
        iou = calculate_iou(predicted_box, ground_truth_box)
        iou_scores.append(iou)
        print(f"      Image: {os.path.basename(test_img_path)} | IoU Score: {iou:.4f}")

if iou_scores:
    average_iou = sum(iou_scores) / len(iou_scores)
    end_time = time.time()
    print("\nRESULTS for '{product}':")
    print(f"  Average IoU: {average_iou:.4f}")
    print(f"  Time Taken: {end_time - start_time:.2f} seconds")
else:
    print(f"  No other images were available to test for product '{product}'")
    print("-" * 50)
else:
    print(f"FATAL ERROR: The dataset path does not exist: {dataset_path}")

```

→ Installing build dependencies ... done
 Getting requirements to build wheel ... done
 Preparing metadata (pyproject.toml) ... done
 `sam2_extract.py` created successfully.
 Starting the main object tracking script...
 --- File Discovery Debug Table ---

Product	Images	Masks
can_chowder	50	75
can_soymilk	50	75
can_tomatosoup	50	75
carton_oj	50	75
carton_soymilk	50	75
diet_coke	50	75
hc_potroastsoup	50	75
juicebox	50	75
rice_tuscan	50	75
ricepilaf	50	75

Processing Product: can_chowder
 Skipping 'can_chowder': Mismatch or not enough files (found 50 images, 75 masks).

```

Processing Product: juicebox
  Skipping 'juicebox': Mismatch or not enough files (found 50 images, 75 masks).
-----
Processing Product: ricepilaf
  Skipping 'ricepilaf': Mismatch or not enough files (found 50 images, 75 masks).
-----
Processing Product: hc_potroastsoup
  Skipping 'hc_potroastsoup': Mismatch or not enough files (found 50 images, 75 masks).
-----
Processing Product: rice_tuscan
  Skipping 'rice_tuscan': Mismatch or not enough files (found 50 images, 75 masks).
-----
Processing Product: carton_oj
  Skipping 'carton_oj': Mismatch or not enough files (found 50 images, 75 masks).
-----
Processing Product: can_soymilk
  Skipping 'can_soymilk': Mismatch or not enough files (found 50 images, 75 masks).
-----
Processing Product: can_tomatosoup
  Skipping 'can_tomatosoup': Mismatch or not enough files (found 50 images, 75 masks).
-----
Processing Product: diet_coke
  Skipping 'diet_coke': Mismatch or not enough files (found 50 images, 75 masks).
-----
Processing Product: carton_soymilk
  Skipping 'carton_soymilk': Mismatch or not enough files (found 50 images, 75 masks).
-----
```

```

import os
import numpy as np
from PIL import Image
import matplotlib.pyplot as plt
import time
from collections import defaultdict

# --- Step 1: Ensure sam2_extract.py is created correctly ---
# This code is reused from before and is correct.
sam2_code = """
import torch
from sam2.build_sam import build_sam2_video_predictor
from PIL import Image
import numpy as np
import matplotlib.pyplot as plt
import os, glob, shutil
import matplotlib.patches as patches

checkpoint = "/content/sam2_hiera_tiny.pt"
model_cfg = "sam2_hiera_t.yaml"

if not os.path.exists(checkpoint):
    raise FileNotFoundError(f"CRITICAL ERROR: Model checkpoint not found at: {checkpoint}")

predictor_vid = build_sam2_video_predictor(model_cfg, checkpoint, device='cuda')
tempfolder = "./tempdir"

def create_if_not_exists(dirname):
    if not os.path.exists(dirname):
        os.mkdir(dirname)

def cleardir(tempfolder):
    if os.path.exists(tempfolder):
        for filepath in glob.glob(tempfolder + "/*"):
            os.unlink(filepath)

def show_mask(mask, ax, obj_id=None):
    cmap = plt.get_cmap("tab10")
    color = np.array([*cmap(obj_id or 0)[:3], 0.6])
```

```

h, w = mask.shape[-2:]
mask_image = mask.reshape(h, w, 1) * color.reshape(1, 1, -1)
ax.imshow(mask_image)

def track_item_boxes(imgpath1, imgpath2, img1boxclasslist, visualize=True):
    create_if_not_exists(tempfolder)
    cleardir(tempfolder)
    shutil.copy(imgpath1, os.path.join(tempfolder, "00000.jpg"))
    shutil.copy(imgpath2, os.path.join(tempfolder, "00001.jpg"))
    inference_state = predictor_vid.init_state(video_path="./tempdir")
    ann_frame_idx = 0
    for img1boxclass in img1boxclasslist:
        ([xmin, xmax, ymin, ymax], objectnumint) = img1boxclass
        box = np.array([xmin, ymin, xmax, ymax], dtype=np.float32)
        predictor_vid.add_new_points_or_box(
            inference_state=inference_state,
            frame_idx=ann_frame_idx,
            obj_id=objectnumint,
            box=box,
        )
    video_segments = {}
    for out_frame_idx, out_obj_ids, out_mask_logits in predictor_vid.propagate_in_video(inference_state):
        video_segments[out_frame_idx] = {
            out_obj_id: (out_mask_logits[i] > 0.0).cpu().numpy()
            for i, out_obj_id in enumerate(out_obj_ids)
        }
    if visualize:
        print("\n--- Visualization ---")
        ([xmin, xmax, ymin, ymax], _) = img1boxclasslist[0]
        fig, axes = plt.subplots(1, 2, figsize=(12, 6))

        # Display the source image with the bounding box
        axes[0].set_title(f"Input: Bounding Box on First Image")
        axes[0].imshow(Image.open(os.path.join(tempfolder, "00000.jpg")))
        rect = patches.Rectangle((xmin, ymin), xmax - xmin, ymax - ymin, linewidth=2, edgecolor='r', facecolor='none')
        axes[0].add_patch(rect)
        axes[0].axis('off')

        # Display the target image with the predicted mask
        axes[1].set_title(f"Output: Predicted Mask on Second Image")
        axes[1].imshow(Image.open(os.path.join(tempfolder, "00001.jpg")))
        if 1 in video_segments and 1 in video_segments[1]:
            for out_obj_id, out_mask in video_segments[1].items():
                show_mask(out_mask, axes[1], obj_id=out_obj_id)
        axes[1].axis('off')
        plt.tight_layout()
        plt.show()

    return video_segments
"""

with open("sam2_extract.py", "w") as f:
    f.write(sam2_code)
print(`sam2_extract.py` created successfully.)
from sam2_extract import track_item_boxes

# --- Step 2: Main processing script with corrected logic and visualization ---

def process_img_png_mask(image_path, mask_path):
    # This helper function remains the same
    try:
        with Image.open(mask_path) as mask_img:
            mask = np.array(mask_img.convert('L')) > 0
            rows, cols = np.any(mask, axis=1), np.any(mask, axis=0)
            if not np.any(rows) or not np.any(cols): return None
            ymin, ymax = np.where(rows)[0][[0, -1]]

```

```

xmin, xmax = np.where(cols)[0][[0, -1]]
return [xmin, xmax, ymin, ymax]
except (FileNotFoundException, Exception):
    return None

print("Starting the main object tracking script...")
dataset_path = '/content/CMU10_3D/data_2D/'

if os.path.exists(dataset_path):
    all_files = os.listdir(dataset_path)

# --- NEW, CORRECTED FILE MATCHING LOGIC ---
# We will find pairs of (image, mask) instead of grouping them separately.
image_mask_pairs = []
all_masks = [f for f in all_files if f.endswith('.png')]

for img_file in sorted([f for f in all_files if f.endswith('.jpg')]):
    # The base name of the image, e.g., 'can_chowder_000001'
    base_name = os.path.splitext(img_file)[0]

    # Find the first mask file that starts with the same base name
    found_mask = next((m for m in all_masks if m.startswith(base_name)), None)

    if found_mask:
        image_mask_pairs.append({
            "product": '_'.join(base_name.split('_')[:-1]),
            "image": os.path.join(dataset_path, img_file),
            "mask": os.path.join(dataset_path, found_mask)
        })

if not image_mask_pairs:
    print("Could not find any valid image-mask pairs. Please check file names.")
else:
    print(f"Found {len(image_mask_pairs)} valid image-mask pairs.")

# --- Let's demonstrate with ONE example for visualization ---
print("Running visualization for one example...")

# Use the first two pairs from the list for tracking
first_pair = image_mask_pairs[0]
second_pair = image_mask_pairs[1]

# Get the bounding box from the very first image's mask
initial_bbox = process_img_png_mask(first_pair["image"], first_pair["mask"])

if initial_bbox:
    # Track the object from the first image to the second image
    # Set visualize=True to see the output images!
    track_item_boxes(
        first_pair["image"],
        second_pair["image"],
        [(initial_bbox, 1)],
        visualize=True
    )
else:
    print("Could not get an initial bounding box for the first image.")
else:
    print(f"FATAL ERROR: The dataset path does not exist: {dataset_path}")

```

```
→ `sam2_extract.py` created successfully.
Starting the main object tracking script...
Found 500 valid image-mask pairs.
Running visualization for one example...
frame loading (JPEG): 100%|██████████| 2/2 [00:00<00:00, 22.39it/s]
/usr/local/lib/python3.11/dist-packages/sam2/sam2_video_predictor.py:786: UserWarning: /usr/local/lib/python3.11/d

Skipping the post-processing step due to the error above. You can still use SAM 2 and it's OK to ignore the error
pred_masks_gpu = fill_holes_in_mask_scores(
propagate in video: 100%|██████████| 2/2 [00:00<00:00, 7.79it/s]
```

Original Image Object



Detected Object in Test Image



```
import os
import numpy as np
from PIL import Image
import matplotlib.pyplot as plt
import matplotlib.patches as patches
from collections import defaultdict

# Make sure to import the tracking function from your script
from sam2_extract import track_item_boxes
```

```
# --- Helper Functions (Copied from your previous script) ---
```

<https://colab.research.google.com/drive/1uOuJNVcBYxv4tC9DS4VsRBpAwewkv5K9#scrollTo=oeHyIOOM23yb&printMode=true>

```

def process_img_png_mask(image_path, mask_path):
    try:
        with Image.open(mask_path) as mask_img:
            mask = np.array(mask_img.convert('L')) > 0
        rows, cols = np.any(mask, axis=1), np.any(mask, axis=0)
        if not np.any(rows) or not np.any(cols): return None
        ymin, ymax = np.where(rows)[0][[0, -1]]
        xmin, xmax = np.where(cols)[0][[0, -1]]
        return [xmin, ymin, xmax, ymax] # Using (xmin, ymin, xmax, ymax) for clarity
    except (FileNotFoundException, Exception):
        return None

def get_bbox_from_mask(mask):
    rows, cols = np.any(mask, axis=1), np.any(mask, axis=0)
    if not np.any(rows) or not np.any(cols): return [0, 0, 0, 0]
    ymin, ymax = np.where(rows)[0][[0, -1]]
    xmin, xmax = np.where(cols)[0][[0, -1]]
    return [xmin, ymin, xmax, ymax]

def calculate_iou(boxA, boxB):
    xA = max(boxA[0], boxB[0]); yA = max(boxA[1], boxB[1])
    xB = min(boxA[2], boxB[2]); yB = min(boxA[3], boxB[3])
    interArea = max(0, xB - xA) * max(0, yB - yA)
    boxAArea = (boxA[2] - boxA[0]) * (boxA[3] - boxA[1])
    boxBArea = (boxB[2] - boxB[0]) * (boxB[3] - boxB[1])
    iou = interArea / float(boxAArea + boxBArea - interArea)
    return iou

# --- Main Debugging Script ---
print("🚀 Starting debug script...")
dataset_path = '/content/CMU10_3D/data_2D'
all_files = os.listdir(dataset_path)
all_masks = [f for f in all_files if f.endswith('.png')]

# --- Find a valid pair to test ---
first_image_path, first_mask_path = None, None
test_image_path, test_mask_path = None, None

for img_file in sorted(f for f in all_files if f.startswith('can_chowder') and f.endswith('.jpg')):
    base_name = os.path.splitext(img_file)[0]
    found_mask = next((m for m in all_masks if m.startswith(base_name)), None)

    if found_mask:
        if first_image_path is None:
            first_image_path = os.path.join(dataset_path, img_file)
            first_mask_path = os.path.join(dataset_path, found_mask)
        else:
            test_image_path = os.path.join(dataset_path, img_file)
            test_mask_path = os.path.join(dataset_path, found_mask)
            break # We have found our two images

# --- Run the detection and visualization ---
if first_image_path and test_image_path:
    print(f"Setup Image: {os.path.basename(first_image_path)}")
    print(f"Test Image: {os.path.basename(test_image_path)}")

    # 1. Get the initial box from the first image
    initial_bbox = process_img_png_mask(first_image_path, first_mask_path)

    # 2. Get the ground truth box for the test image
    ground_truth_box = process_img_png_mask(test_image_path, test_mask_path)

    # 3. Get the predicted box from the model
    video_segments = track_item_boxes(first_image_path, test_image_path, [[initial_bbox[0], initial_bbox[2], initial_bt
predicted_box = [0, 0, 0, 0]
if 1 in video_segments and 1 in video_segments[1].

```

```
    + in video_segments and in video_segments[1].  
predicted_box = get_bbox_from_mask(video_segments[1][1])  
  
# 4. Print the coordinates  
print(f"\nGround Truth Box: {ground_truth_box}")  
print(f"Predicted Box:      {predicted_box}")  
  
# 5. Calculate and print the IoU  
iou = calculate_iou(ground_truth_box, predicted_box)  
print(f"\nCalculated IoU: {iou:.4f}")  
  
# 6. Visualize the result  
fig, ax = plt.subplots(1, 1, figsize=(8, 8))  
ax.imshow(Image.open(test_image_path))  
ax.set_title("Detection Result")  
  
# Draw Ground Truth (Green)  
if ground_truth_box:  
    (xmin, ymin, xmax, ymax) = ground_truth_box  
    gt_rect = patches.Rectangle((xmin, ymin), xmax - xmin, ymax - ymin, linewidth=2, edgecolor='g', facecolor='none')  
    ax.add_patch(gt_rect)  
  
# Draw Prediction (Red)  
if predicted_box != [0, 0, 0, 0]:  
    (xmin, ymin, xmax, ymax) = predicted_box  
    pred_rect = patches.Rectangle((xmin, ymin), xmax - xmin, ymax - ymin, linewidth=2, edgecolor='r', facecolor='none')  
    ax.add_patch(pred_rect)  
  
ax.legend()  
plt.show()  
  
else:  
    print("Could not find two 'can_chowder' images to test.")
```

```
→ Starting debug script...
Setup Image: can_chowder_000001.jpg
Test Image: can_chowder_000002.jpg
frame loading (JPEG): 100%|██████████| 2/2 [00:00<00:00, 31.76it/s]
propagate in video: 100%|██████████| 2/2 [00:00<00:00,  9.61it/s]
\nGround Truth Box: [np.int64(249), np.int64(148), np.int64(336), np.int64(278)]
Predicted Box:      [np.int64(150), np.int64(0), np.int64(278), np.int64(0)]
\nCalculated IoU: 0.0000
```

Detection Result



Starting debug script...

Setup Image: can_chowder_000001.jpg

Test Image: can_chowder_000002.jpg

frame loading (JPEG): 100%|██████████| 2/2 [00:00<00:00, 31.76it/s]

propagate in video: 100%|██████████| 2/2 [00:00<00:00, 9.61it/s]

\nGround Truth Box: [np.int64(249), np.int64(148), np.int64(336), np.int64(278)]

Predicted Box: [np.int64(150), np.int64(0), np.int64(278), np.int64(0)]

\nCalculated IoU: 0.0000

Detection Result

